

# Week 1 — Cloud Computing as a Systems Discipline

## Foundations: Scale, Failure, Declarative Control

Dr. J. Burns

INFT 3611

November 29, 2025

### **Section 1: Introduction**

*This content is based on the following public resources: <https://www.distributed-systems.net/index.php/books/ds4/>*

# Cloud Computing as a Systems Discipline

## Core Perspective

Cloud computing is fundamentally a **distributed systems discipline**, where correctness depends on algorithmic coordination across unreliable and asynchronous components.

- Cloud platforms behave more like planet-scale operating systems than virtualized datacenters [1].
- Core constraints include:
  - absence of globally consistent state,
  - nondeterministic event ordering,
  - continuous node and process churn.
- These conditions require distributed algorithms that ensure convergence despite message delay and partial observability.

# Why Cloud Requires Systems Thinking

## Misconception

“Cloud computing is just servers in someone else’s datacenter.”

- Large production systems (Google, AWS, Azure) operate under **persistent instability**: crashes, partitions, and restarts occur continuously [2].
- Procedural logic assumes:
  - linear execution,
  - reliable infrastructure,
  - timely intervention.These assumptions fail at scale.
- Correctness requires:
  - idempotent operations,
  - eventual consistency,
  - monotonic state evolution [3].

# Cloud as a Large-Scale Operating System

## Analogy

Cloud control planes approximate operating system kernels, but operate across unreliable devices, asynchronous communication, and weak consistency.

- Responsibilities include:
  - global workload scheduling [4],
  - policy governance,
  - identity and isolation,
  - reconciliation of desired and observed state.
- Unlike local kernels, cloud systems cannot rely on:
  - atomic memory,
  - global monotonic time,
  - reliable message delivery.
- Correctness emerges from distributed convergence, not sequential execution.

# Design Tensions in Cloud Architecture

## Structural Tensions

Global-scale systems must balance fundamental constraints that cannot be optimized simultaneously.

- **Consistency vs throughput:** strong consistency requires coordination, limiting parallelism [5].
- **Elasticity vs predictability:** scaling shifts resource topology unpredictably.
- **Automation vs visibility:** automation improves correctness but reduces operator observability.
- **Availability vs correctness:** high availability often permits temporary inconsistency.

# Scale Amplifies Complexity

## Scaling Reality

At cloud scale, rare events become constant. Small failure probabilities multiplied across millions of components produce continuous disturbances.

- Industry data shows that large fleets experience ongoing hardware and software faults [6].
- The system behaves like a high-dimensional dynamical system under continuous perturbation.
- Manual reasoning cannot track or correct drift at this speed.
- Stability requires automated control loops that re-establish invariants despite ongoing disturbances.

# Continuous Perturbation as the Normal State

## Operational Reality

At large scale, the system is **never quiescent**. Failures, recoveries, and timing anomalies form a continuous background process.

- Node churn (joins, leaves, crashes) occurs constantly in large fleets [2].
- Network delay, jitter, and loss vary dynamically across datacenter fabrics.
- Clock skew accumulates between nodes, undermining assumptions about event ordering [3].
- Replicated components diverge due to nondeterministic propagation delays.
- Systems must rely on convergence, not instantaneous correctness.

# The State Explosion Problem

## Key Idea

Cloud systems maintain multiple interacting state domains whose cross-product produces an intractably large global state space.

- State classes include:
  - configuration state,
  - runtime and health state,
  - dependency and topology state,
  - policy and security state.
- Even small inconsistencies can cascade across these interconnected layers.
- Strongly consistent global snapshots require coordination mechanisms such as consensus [3], which becomes expensive at scale.
- Automated controllers are required to continually manage drift within this state space.

# Limits of Manual Coordination

## Human Limitation

Human operators cannot observe, reason about, or repair distributed state quickly enough to maintain correctness in highly concurrent environments.

- Manual intervention introduces variable latency, allowing inconsistencies to widen.
- Operators often apply partial or inconsistent fixes, unintentionally increasing drift.
- Post-incident reports frequently attribute outages to configuration errors and slow reaction paths [6].
- The rate of topological change exceeds human cognitive limits.

# Why Automation Becomes Mandatory

## Benefits of Automated Controllers

Automation enforces system invariants at speeds and levels of consistency impossible for human operators.

- Automated controllers:
  - detect drift immediately,
  - execute corrective actions uniformly,
  - operate continuously rather than episodically.
- Organizations such as Google and AWS rely on continuous automated monitoring and reconciliation [6].
- Automation transforms correctness from manual labor into a property maintained by algorithmic control.

# Coordination as a Control Problem

## Control-Theoretic Interpretation

Distributed systems resemble feedback control systems, where noisy observations and delayed actions must still lead to stable behavior.

- The core loop: **Observe → Compare → Correct**.
- Controllers must maintain stability under:
  - partial observability,
  - asynchronous updates,
  - unpredictable disturbances,
  - variable execution latency.
- The primary correctness property is **convergence**, not perfect instantaneous agreement.
- This insight underpins modern control-plane architectures used in Kubernetes, Borg, and distributed service fabrics.

# Failure as a First-Class Property

## Principle

In cloud environments, **failure is a continuous operating condition**, not an exceptional event.

- Large-scale systems observe constant hardware faults, process exits, link failures, and retries [2].
- Components fail independently, producing diverse and unpredictable patterns.
- Software must assume that any communication, computation, or storage operation may not complete as expected.
- Reliability emerges from redundancy, feedback loops, and corrective algorithms—never from individual component durability.

## Fundamental Constraint

Distributed systems never possess a fully accurate, real-time global view of state.

- Observations are:
  - delayed,
  - partial,
  - stale,
  - inconsistent across replicas.
- Controllers must take action despite incomplete knowledge.
- Event timestamps cannot be fully trusted due to clock skew and clock drift [3].
- Systems must be engineered to remain stable under uncertainty.

## Common Resilience Techniques

Distributed systems use operational properties that remain correct under retries, reordering, and partial execution.

- **Idempotent operations:** issuing the same update multiple times produces the same final state.
- **Monotonic state transitions:** state evolves in one direction, tolerating out-of-order messages.
- **Eventual convergence:** all replicas converge on a stable state after disturbances.
- These principles strengthen correctness in the presence of nondeterministic failures and message delivery conditions.

# Why Declarative Instead of Procedural

## Issue

Procedural instructions rely on ordering guarantees that distributed systems cannot provide.

- A sequence of procedural steps will fail if:
  - messages are delayed,
  - node scheduling varies,
  - operations are retried,
  - partial failures interrupt execution.
- Declarative intent defines the target invariant; controllers compute the needed transitions.
- This shifts correctness from step-by-step execution toward long-run convergence [3].

# Intent as System Specification

## Definition

**Declarative intent** specifies the desired stable properties of the system rather than the sequence of steps to achieve them.

- Operators express targets such as:
  - replica counts,
  - topology constraints,
  - policy and security requirements,
  - resource limits.
- Controllers continuously reconcile observed state with this specification.
- Self-healing behavior emerges because the intent remains stable even as the environment changes.

## Sources of Drift

Differences between an operator's declared intent and the system's actual condition accumulate continuously in distributed environments.

- Causes of drift:
  - node failures and evictions,
  - message delays and reordering,
  - resource contention,
  - propagation latency between replicas.
- Drift is unavoidable; systems must detect and correct it rather than assume stable state.
- Persistent automated reconciliation is required to maintain correctness at scale.

# Desired State

## Definition

The **desired state** is the operator-defined specification of how the system should behave under correct conditions.

- Examples:
  - “Maintain 5 replicas of service X.”
  - “Ensure all nodes enforce network policy Y.”
  - “Keep workloads within zone and quota constraints.”
- The desired state is stable, declarative, and often version-controlled.
- It acts as the reference point for control-plane decisions.

# Observed State

## Reality

The **observed state** reflects the system's actual behavior but is always partial, stale, and approximate in distributed environments.

- Observed state is influenced by:
  - network delays,
  - local caching,
  - clock skew,
  - inconsistent updates across replicas.
- No controller sees the whole cluster instantaneously.
- Decision making must tolerate incomplete, outdated, or ambiguous observations [3].

# State Comparison

## Mechanism

Control planes continuously compute corrective actions based on the difference:

$$\text{Correction} = \text{Desired State} - \text{Observed State}$$

- Controllers identify:
  - missing resources (under-replication),
  - excess resources (over-provisioning),
  - inconsistent policy or configuration.
- This differential approach is robust to reordering and partial failure.
- Reconciliation logic is idempotent and monotonic, enabling safe repeated execution.

# Reconciliation Loop

## Control Loop Model

The canonical control-plane loop:

**Observe** → **Compare** → **Correct**

drives convergence toward the declared intent.

- Controllers re-evaluate state continuously, not on demand.
- Reconciliation absorbs disturbances caused by failures, delays, and nondeterministic event ordering.
- Long-term correctness is achieved through **eventual convergence**, not perfect instantaneous agreement [6].
- This strategy underlies modern systems such as Kubernetes, Borg, and Mesos.

# Idempotence and Safety

## Key Principle

Corrective operations must be **idempotent**: safe to repeat, reorder, or retry without producing unintended side effects.

- Idempotence protects correctness when:
  - messages are duplicated,
  - operations are retried,
  - partial failures interrupt updates,
  - controllers race to modify shared state.
- Modern control-plane implementations, including Borg and Kubernetes, rely heavily on idempotent handlers for stable reconciliation [4].
- Idempotence forms the “safety net” for distributed convergence.

# Distributed Reconciliation Dynamics

## Challenges in Multi-Controller Systems

Multiple controllers operate concurrently and may interact in unexpected ways.

- Risks include:
  - oscillation (controllers undo each other's changes),
  - race conditions in shared state,
  - inconsistent enforcement of cross-cutting policies.
- Stability requires:
  - clearly defined ownership of state domains,
  - deterministic conflict-resolution strategies,
  - eventually consistent propagation of updates.
- Formal reasoning increasingly informs the design of modern controllers [5].

# Case Study: Replica Restoration

## Scenario

A replicated service must maintain the invariant: **exactly  $N$  healthy replicas** should exist.

- Failure events remove replicas (crashes, eviction, unreachable nodes).
- The observed state becomes: actual replicas  $< N$ .
- Controllers perform:
  - detection of missing replicas,
  - scheduling or placement decisions,
  - instantiation of new replicas.
- Even under partial failure, the system converges back to the declared invariant through automated reconciliation.

# Case Study: Policy Enforcement Under Failure

## Scenario

A network or security policy must be enforced globally, but partial failures lead to inconsistent rule deployment.

- Individual nodes may:
  - fail to apply required rules,
  - apply outdated or conflicting rules,
  - lose local state due to restart or eviction.
- Controllers detect drift via periodic inspection and update propagation.
- Convergence resembles epidemic or anti-entropy protocols used in distributed databases [5].
- The system gradually restores global policy consistency without manual intervention.

# Case Study: Elastic Scaling

## Scenario

Workload load increases or decreases, triggering an elasticity policy governed by declarative constraints.

- The desired state encodes elasticity bounds (min/max replicas, scaling rules).
- The observed state reflects resource saturation or underutilization.
- Controllers:
  - resize replica sets,
  - redistribute load,
  - rebalance across zones or failure domains.
- Elastic scaling loops must avoid oscillation and ensure stable response dynamics [6].

# Self-Stabilizing Control Planes

## Insight

Modern cloud control planes approximate **self-stabilizing distributed systems** that converge toward correctness from arbitrary states.

- The concept originates from Dijkstra's theory of self-stabilization.
- Control planes repeatedly correct drift through reconciliation loops.
- Convergence occurs despite:
  - stale data,
  - transient inconsistency,
  - partial failures,
  - message reordering.
- This explains why declarative cloud systems recover automatically from widespread disturbances.

# Why Declarative Control Dominates

## Architectural Reasoning

Declarative control is the only scalable strategy for correctness in asynchronous, failure-prone environments.

- Step-by-step procedural logic breaks under concurrency, races, and retries.
- Declarative intent provides a stable “truth anchor” for reconciliation.
- Long-run correctness derives from the convergence of multiple controllers acting independently.
- This pattern underlies the design of Kubernetes, Borg, Mesos, and large cloud fabrics [4].

# Summary of Week 1

## Key Takeaways

- Cloud computing must be analyzed as a distributed systems discipline.
- Scale introduces continuous failure, partial observability, and nondeterminism.
- Declarative intent defines system invariants; observed state reflects imperfect reality.
- Reconciliation (Observe → Compare → Correct) is the core mechanism for achieving convergence.
- Idempotence, monotonicity, and self-stabilization ensure robustness under reordering and retries.

# Suggested Activities

## Activities for Students

- Analyze a distributed system (e.g., Cassandra, etcd) and identify how it handles drift.
- Explain how idempotence prevents race conditions and supports retries.
- Compare reconciliation-based coordination vs consensus-based coordination.
- Review an incident report and identify how declarative or procedural control contributed to system stability or instability.

Questions?  
Discussion?

Preparation for Week 2: Models of Distributed Computation

## References I

-  A. Verma, L. Cherkasova, R. Campbell, and G. Shen, "Large-scale cluster management at Google," *Proceedings of the 2008 International Conference on Autonomic Computing*, 2008.
-  J. Hamilton, "On Designing and Deploying Internet-Scale Services," *Proceedings of the Large Installation System Administration Conference (LISA)*, 2007.
-  L. Lamport, "Time, Clocks, and the Ordering of Events in a Distributed System," *Communications of the ACM*, vol. 21, no. 7, pp. 558–565, 1978.
-  A. Verma et al., "Large-scale cluster management at Google with Borg," *Proceedings of the European Conference on Computer Systems (EuroSys)*, 2015.
-  S. Gilbert and N. Lynch, "Brewer's conjecture and the feasibility of consistent, available, partition-tolerant web services," *ACM SIGACT News*, vol. 33, no. 2, pp. 51–59, 2002.
-  B. Beyer, C. Jones, J. Petoff, and N. R. Murphy, *Site Reliability Engineering: How Google Runs Production Systems*. O'Reilly Media, 2016.

- A. Verma, L. Cherkasova, R. Campbell, and G. Shen, “Large-scale cluster management at google,” in *International Conference on Autonomic Computing*, 2008.
- J. Hamilton, “On designing and deploying internet-scale services,” *LISA*, 2007.
- L. Lamport, “Time, clocks, and the ordering of events in a distributed system,” *Communications of the ACM*, vol. 21, no. 7, pp. 558–565, 1978.
- A. Verma *et al.*, “Large-scale cluster management at google with borg,” in *EuroSys*, 2015.
- S. Gilbert and N. Lynch, “Brewer’s conjecture and the feasibility of consistent, available, partition-tolerant web services,” *SIGACT News*, vol. 33, no. 2, pp. 51–59, 2002.
- B. Beyer, C. Jones, J. Petoff, and N. R. Murphy, *Site Reliability Engineering*. O’Reilly Media, 2016.